

April 25, 2023

- 1 Group Name: Go Bear!
- 2 Name: Xiaoke Song
- 3 Email: xiaokesong57@gmail.com
- 4 Country: born in China, college in the US
- 5 College: UC Berkeley
- 6 Specialization: Data Science

```
[1]: import pandas as pd
```

```
[2]: df = pd.read_csv("cust_seg.csv")
      df.head(5)
```

```
/var/folders/_0/nmpfpzw134n12j0c0z6jtrw80000gn/T/ipykernel_5258/1520097819.py:1:
DtypeWarning: Columns (16) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
df = pd.read_csv("cust_seg.csv")
```

```
[2]: Unnamed: 0  fecha_dato  ncodpers  ind_empleado  pais_residencia  sexo  age  \
0            0  2015-01-28   1375586             N             ES    H    35
1            1  2015-01-28   1050611             N             ES    V    23
2            2  2015-01-28   1050612             N             ES    V    23
3            3  2015-01-28   1050613             N             ES    H    22
4            4  2015-01-28   1050614             N             ES    V    23
```

```
      fecha_alta  ind_nuevo  antiguedad  ...  ind_hip_fin_ult1  ind_plan_fin_ult1  \
0  2015-01-12         0.0           6  ...             0             0
1  2012-08-10         0.0          35  ...             0             0
2  2012-08-10         0.0          35  ...             0             0
3  2012-08-10         0.0          35  ...             0             0
4  2012-08-10         0.0          35  ...             0             0
```

```
      ind_pres_fin_ult1  ind_reca_fin_ult1  ind_tjcr_fin_ult1  ind_valo_fin_ult1  \
```

0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

	ind_viv_fin_ult1	ind_nomina_ult1	ind_nom_pens_ult1	ind_recibo_ult1
0	0	0.0	0.0	0
1	0	0.0	0.0	0
2	0	0.0	0.0	0
3	0	0.0	0.0	0
4	0	0.0	0.0	0

[5 rows x 48 columns]

Problem description: Customer Segmentation \_\_\_\_ “XYZ bank wants to roll out Christmas offers to their customers. But Bank does not want to roll out same offer to all customers instead they want to roll out personalized offer to particular set of customers. If they manually start understanding the category of customer then this will be not efficient and also they will not be able to uncover the hidden pattern in the data ( pattern which group certain kind of customer in one category). Bank approached ABC analytics company to solve their problem. Bank also shared information with ABC analytics that they don’t want more than 5 group as this will be inefficient for their campaign.”

Data Understanding: The data contains information about customers of the bank, including demographic information, customer behavior, and product ownership. There are several columns in the dataset, including customer code, age, seniority, activity index, gross income, and various product ownership indicators.

What type of data you have got for analysis? — Based on the provided information, we have a dataset containing customer information for XYZ bank. The dataset includes demographic information, customer behavior, and product ownership, and it appears to be in a structured format. The warning message “Columns (16) have mixed types” indicates that there might be some columns with mixed data types, but without further information, it’s unclear which columns are affected and what the specific data types are. Overall, it seems like we have tabular data that can be analyzed using various techniques such as descriptive statistics, data visualization, and clustering algorithms for customer segmentation.

```
[4]: # check for missing values
print('Number of missing values in each column:')
print(df.isnull().sum())
```

```
Number of missing values in each column:
Unnamed: 0          0
fecha_dato          0
ncodpers            0
ind_empleado       10782
pais_residencia    10782
sexo               10786
```

age	0
fecha_alta	10782
ind_nuevo	10782
antiguedad	0
indrel	10782
ult_fec_cli_1t	998899
indrel_1mes	10782
tiprel_1mes	10782
indresi	10782
indext	10782
conyuemp	999822
canal_entrada	10861
indfall	10782
tipodom	10782
cod_prov	17734
nomprov	17734
ind_actividad_cliente	10782
renta	175183
ind_ahor_fin_ult1	0
ind_aval_fin_ult1	0
ind_cco_fin_ult1	0
ind_cder_fin_ult1	0
ind_cno_fin_ult1	0
ind_ctju_fin_ult1	0
ind_ctma_fin_ult1	0
ind_ctop_fin_ult1	0
ind_ctpp_fin_ult1	0
ind_deco_fin_ult1	0
ind_deme_fin_ult1	0
ind_dela_fin_ult1	0
ind_ecue_fin_ult1	0
ind_fond_fin_ult1	0
ind_hip_fin_ult1	0
ind_plan_fin_ult1	0
ind_pres_fin_ult1	0
ind_reca_fin_ult1	0
ind_tjcr_fin_ult1	0
ind_valo_fin_ult1	0
ind_viv_fin_ult1	0
ind_nomina_ult1	5402
ind_nom_pens_ult1	5402
ind_recibo_ult1	0

dtype: int64

There are missing values in the following columns:

ind\_empleado,  
pais\_residencia,

```

sexo,
fecha_alta,
ind_nuevo,
indrel,
ult_fec_cli_1t,
indrel_1mes,
tiprel_1mes,
indresi,
indext,
conyuemp,
canal_entrada,
indfall,
tipodom,
cod_prov,
nomprov,
ind_actividad_cliente,
renta,
ind_nomina_ult1,
ind_nom_pens_ult1;

```

The column “ult\_fec\_cli\_1t” has a very large number of missing values, while the columns “ind\_empleado”, “pais\_residencia”, “sexo”, “fecha\_alta”, “ind\_nuevo”, “indrel”, “indrel\_1mes”, “tiprel\_1mes”, “indresi”, “indext”, “conyuemp”, “canal\_entrada”, “indfall”, “tipodom”, “cod\_prov”, “nomprov”, “ind\_actividad\_cliente”, “renta”, “ind\_nomina\_ult1”, and “ind\_nom\_pens\_ult1” have a relatively smaller number of missing values.

```

[3]: # calculate the summary statistics
summary = df.describe()

# print the summary statistics
print(summary)

```

	Unnamed: 0	ncodpers	ind_nuevo	indrel \
count	1000000.000000	1.000000e+06	989218.000000	989218.000000
mean	499999.500000	6.905967e+05	0.000489	1.109074
std	288675.278933	4.044084e+05	0.022114	3.267624
min	0.000000	1.588900e+04	0.000000	1.000000
25%	249999.750000	3.364110e+05	0.000000	1.000000
50%	499999.500000	6.644760e+05	0.000000	1.000000

75%	749999.250000	1.074511e+06	0.000000	1.000000
max	999999.000000	1.379131e+06	1.000000	99.000000

	indrel_1mes	tipodom	cod_prov	ind_actividad_cliente \
count	989218.000000	989218.0	982266.000000	989218.000000
mean	1.000085	1.0	26.852131	0.564971
std	0.012954	0.0	12.422924	0.495761
min	1.000000	1.0	1.000000	0.000000
25%	1.000000	1.0	18.000000	0.000000
50%	1.000000	1.0	28.000000	1.000000
75%	1.000000	1.0	33.000000	1.000000
max	3.000000	1.0	52.000000	1.000000

	renta	ind_ahor_fin_ult1	...	ind_hip_fin_ult1 \
count	8.248170e+05	1000000.000000	...	1000000.000000
mean	1.396462e+05	0.000177	...	0.009982
std	2.389858e+05	0.013303	...	0.099410
min	1.202730e+03	0.000000	...	0.000000
25%	7.157184e+04	0.000000	...	0.000000
50%	1.066519e+05	0.000000	...	0.000000
75%	1.634325e+05	0.000000	...	0.000000
max	2.889440e+07	1.000000	...	1.000000

	ind_plan_fin_ult1	ind_pres_fin_ult1	ind_reca_fin_ult1 \
count	1000000.000000	1000000.000000	1000000.000000
mean	0.014553	0.004661	0.072581
std	0.119755	0.068112	0.259448
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	ind_tjcr_fin_ult1	ind_valo_fin_ult1	ind_viv_fin_ult1 \
count	1000000.000000	1000000.000000	1000000.000000
mean	0.066084	0.039378	0.006442
std	0.248429	0.194493	0.080003
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

	ind_nomina_ult1	ind_nom_pens_ult1	ind_recibo_ult1
count	994598.000000	994598.000000	1000000.000000
mean	0.071629	0.079543	0.166275
std	0.257873	0.270584	0.372327
min	0.000000	0.000000	0.000000

25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000

[8 rows x 33 columns]

The 75th percentile refers to the point in a dataset where 75% of the observations are below that value. Outliers are observations that fall outside of the typical range of values for a dataset, and can be identified using various statistical methods. Common methods for identifying outliers include the use of box plots, z-scores, and interquartile range (IQR).

```
[4]: # calculate the skewness
      skewness = df.skew()

      # print the skewness
      print(skewness)
```

```
Unnamed: 0      -2.511790e-15
ncodpers        3.227844e-02
ind_nuevo       4.517572e+01
indrel          2.992451e+01
indrel_1mes     1.534514e+02
tipodom        0.000000e+00
cod_prov       -1.512119e-01
ind_actividad_cliente -2.621046e-01
renta          5.223413e+01
ind_ahor_fin_ult1 7.514476e+01
ind_aval_fin_ult1 1.601190e+02
ind_cco_fin_ult1 -1.152401e+00
ind_cder_fin_ult1 4.109809e+01
ind_cno_fin_ult1 2.571915e+00
ind_ctju_fin_ult1 8.391620e+00
ind_ctma_fin_ult1 9.903618e+00
ind_ctop_fin_ult1 1.405709e+00
ind_ctpp_fin_ult1 3.309287e+00
ind_deco_fin_ult1 2.145683e+01
ind_deme_fin_ult1 1.773314e+01
ind_dela_fin_ult1 3.467512e+00
ind_ecue_fin_ult1 2.555226e+00
ind_fond_fin_ult1 5.815247e+00
ind_hip_fin_ult1 9.858534e+00
ind_plan_fin_ult1 8.107362e+00
ind_pres_fin_ult1 1.454481e+01
ind_reca_fin_ult1 3.294845e+00
ind_tjcr_fin_ult1 3.493287e+00
ind_valo_fin_ult1 4.736660e+00
ind_viv_fin_ult1 1.233849e+01
```

```
ind_nomina_ult1          3.322353e+00
ind_nom_pens_ult1        3.107782e+00
ind_recibo_ult1          1.792646e+00
dtype: float64
```

```
/var/folders/_0/nmpfpzw134n12j0c0z6jtrw80000gn/T/ipykernel_5258/2832739648.py:2:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    skewness = df.skew()
```

To determine the skewness of the data, we need to look at the signs of the values in the third column. If the value is positive, the distribution is right-skewed (or positively skewed), and if the value is negative, the distribution is left-skewed (or negatively skewed).

From the given data, we can see that:

Unnamed: 0, ncodpers, tipodom, ind\_cco\_fin\_ult1, ind\_cno\_fin\_ult1, ind\_ctop\_fin\_ult1, ind\_ctpp\_fin\_ult1,

ind\_dela\_fin\_ult1, ind\_ecue\_fin\_ult1, ind\_reca\_fin\_ult1, ind\_tjcr\_fin\_ult1, ind\_valo\_fin\_ult1, ind\_nomina\_ult1,

ind\_nom\_pens\_ult1, and ind\_recibo\_ult1 have a skewness value of approximately zero (around 0).

ind\_nuevo, indrel, ind\_cder\_fin\_ult1, ind\_ctju\_fin\_ult1, ind\_ctma\_fin\_ult1, ind\_deco\_fin\_ult1, ind\_deme\_fin\_ult1,

ind\_fond\_fin\_ult1, ind\_hip\_fin\_ult1, ind\_plan\_fin\_ult1, ind\_pres\_fin\_ult1, and ind\_viv\_fin\_ult1 have positive skewness values, indicating a right-skewed distribution. cod\_prov and ind\_actividad\_cliente have negative skewness values, indicating a left-skewed distribution. indrel\_1mes and ind\_aval\_fin\_ult1 have very high positive skewness values, indicating extreme right-skewed distributions. Therefore, we can say that the given dataset has a mix of left-skewed, right-skewed, and extreme right-skewed distributions.

What approaches you are trying to apply on your data set to overcome problems like NA value, outlier etc and why? — one approach to handle missing values is imputation using the mean value of the column. Here's an example code below. In this code, we first load the data and then check for missing values using `isnull()` and `sum()` methods. Then, we replace the missing values with the mean value of the column using `fillna()` method with `inplace=True` to replace the missing values in the original data frame. Finally, we check for missing values again to confirm that all missing values have been replaced with the mean value of the column.

```
[ ]: import numpy as np

# check for missing values
print(df.isnull().sum())

# replace missing values with mean value of column
df.fillna(df.mean(), inplace=True)
```

```
# check for missing values again
print(df.isnull().sum())
```

```
Unnamed: 0          0
fecha_dato          0
ncodpers            0
ind_empleado       10782
pais_residencia    10782
sexo              10786
age                0
fecha_alta         10782
ind_nuevo          10782
antiguedad         0
indrel             10782
ult_fec_cli_1t     998899
indrel_1mes        10782
tiprel_1mes        10782
indresi           10782
indext            10782
conyuemp          999822
canal_entrada      10861
indfall           10782
tipodom           10782
cod_prov          17734
nomprov           17734
ind_actividad_cliente 10782
renta             175183
ind_ahor_fin_ult1  0
ind_aval_fin_ult1  0
ind_cco_fin_ult1   0
ind_cder_fin_ult1  0
ind_cno_fin_ult1   0
ind_ctju_fin_ult1  0
ind_ctma_fin_ult1  0
ind_ctop_fin_ult1  0
ind_ctpp_fin_ult1  0
ind_deco_fin_ult1  0
ind_deme_fin_ult1  0
ind_dela_fin_ult1  0
ind_ecue_fin_ult1  0
ind_fond_fin_ult1  0
ind_hip_fin_ult1   0
ind_plan_fin_ult1  0
ind_pres_fin_ult1  0
ind_reca_fin_ult1  0
ind_tjcr_fin_ult1  0
ind_valo_fin_ult1  0
```



```
ind_viv_fin_ult1      0
ind_nomina_ult1      5402
ind_nom_pens_ult1     5402
ind_recibo_ult1       0
dtype: int64
```

```
[ ]:
```